

Add a New System Call to list all processes

Target

1. Add a new system call into the linux kernel
2. Test the new system call in user model

Tools

Install GCC Software Colletion

```
sudo apt-get install build-essential
```

How to use GCC

- [gcc and make](#)

How to do

see the pdf document: newsyscall_step2.pdf

Step0

new customized kernel config: [config1](#), [config2](#)

Step1 (Linux kernel 5.19)

include/linux/syscalls.h

在文件(No. 1279)

#endif /* CONFIG_ARCH_HAS_SYSCALL_WRAPPER */之前, 添加一行:

```
asmlinkage long sys_schello(void);
```

Step2 (Linux kernel 5.19)

kernel/sys.c

在文件SYSCALL_DEFINE0(gettid)函数之后 (No. 949) , 添加如下行:

```
SYSCALL_DEFINE0(schello)
```

```
{  
    printk("Hello new system call schello!Your ID\n");  
    return 0;  
}
```

Step3 (Linux kernel 5.19)

针对64位OS

arch/x86/entry/syscalls/syscall_64.tbl

在文件334 common memfd_secret sys_memfd_secret 行之后，添加如下行：

```
335 common schello sys_schello
```

Step4

make clean

make -j5

sudo make modules_install

sudo make install

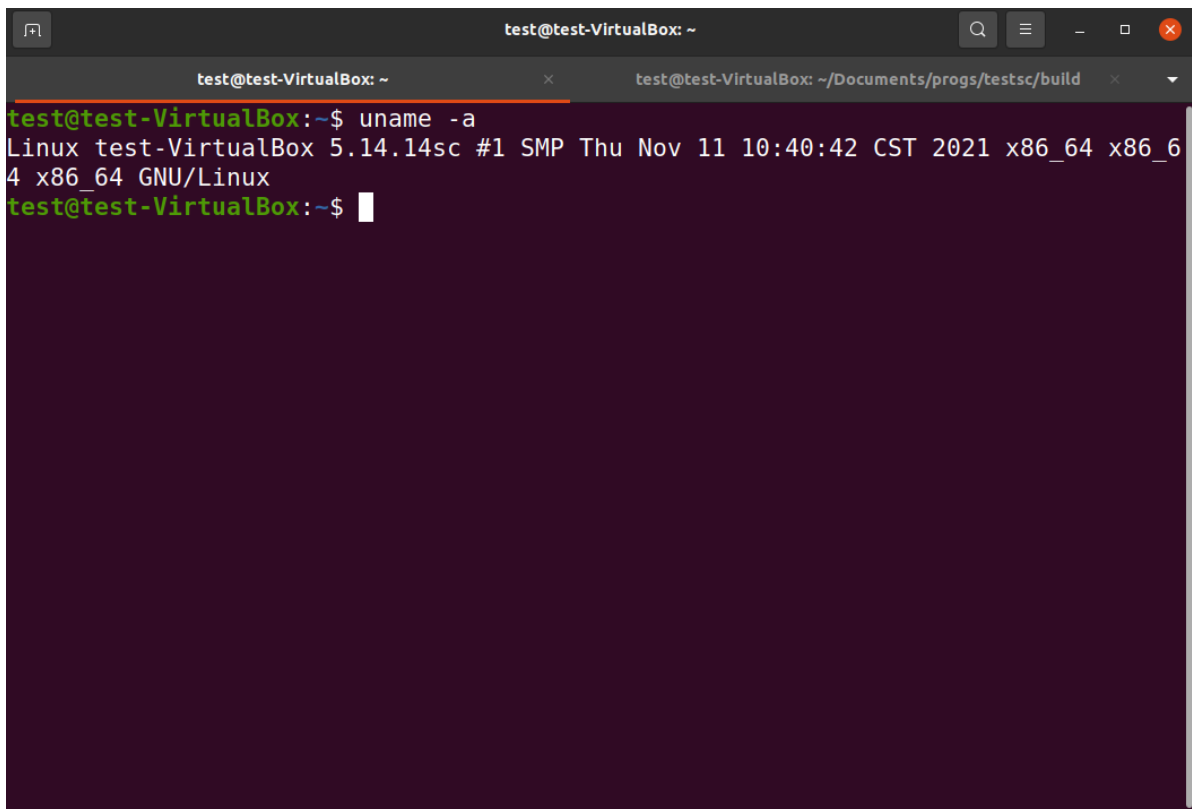
Step 5

重新启动：

reboot

确认新内核是否成功运行：

```
uname -a
```



```
test@test-VirtualBox: ~  
test@test-VirtualBox: ~  
test@test-VirtualBox: ~$ uname -a  
Linux test-VirtualBox 5.14.14sc #1 SMP Thu Nov 11 10:40:42 CST 2021 x86_64 x86_64  
4 x86_64 GNU/Linux  
test@test-VirtualBox:~$
```

Step 6

编写用户态测试程序testschello.c

```
#include <unistd.h>
#include <sys/syscall.h>
#include <sys/types.h>
#include <stdio.h>
#define __NR_schello 336
int main(int argc, char *argv[])
{
    syscall(__NR_schello);
    print("ok! run dmesg | grep hello in terminal!\n");
    return 0;
}
```

Step 7

编译用户态测试程序testschello.c, 并执行

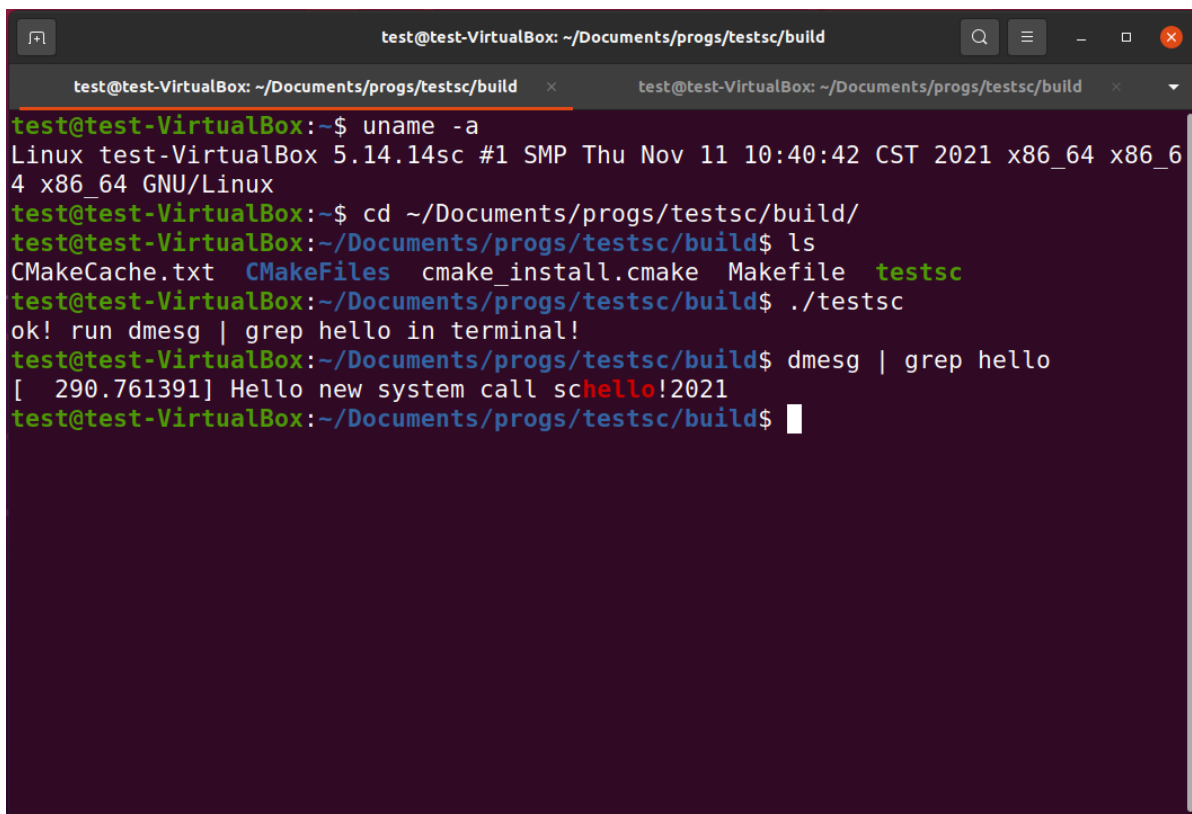
```
gcc -o testsc testschello.c

$ sudo dmesg -C

./testsc

$dmesg | grep schello

[ 1648.215250] Hello new system call schello!
```



```
test@test-VirtualBox: ~/Documents/progs/testsc/build
test@test-VirtualBox:~$ uname -a
Linux test-VirtualBox 5.14.14sc #1 SMP Thu Nov 11 10:40:42 CST 2021 x86_64 x86_64 x86_64 GNU/Linux
test@test-VirtualBox:~$ cd ~/Documents/progs/testsc/build/
test@test-VirtualBox:~/Documents/progs/testsc/build$ ls
CMakeCache.txt  CMakeFiles  cmake_install.cmake  Makefile  testsc
test@test-VirtualBox:~/Documents/progs/testsc/build$ ./testsc
ok! run dmesg | grep hello in terminal!
test@test-VirtualBox:~/Documents/progs/testsc/build$ dmesg | grep hello
[ 290.761391] Hello new system call schello!2021
test@test-VirtualBox:~/Documents/progs/testsc/build$
```

process in Linux kernel

<linux/sched.h>

```
struct task_struct {
#ifdef CONFIG_THREAD_INFO_IN_TASK
    /*
     * For reasons of header soup (see current_thread_info()), this
     * must be the first element of task_struct.
     */
    struct thread_info    thread_info;
#endif
    unsigned int          __state;

#ifdef CONFIG_PREEMPT_RT
    /* saved state for "spinlock sleepers" */
    unsigned int          saved_state;
#endif

    /*
     * This begins the randomizable portion of task_struct. Only
     * scheduling-critical items should be added above here.
     */
    randomized_struct_fields_start

    void                  *stack;
    refcount_t            usage;
    /* Per task flags (PF_*), defined further below: */
    unsigned int          flags;
    unsigned int          ptrace;

#ifdef CONFIG_SMP
    int                   on_cpu;
    struct __call_single_node wake_entry;
    unsigned int          wakee_flips;
    unsigned long          wakee_flip_decay_ts;
    struct task_struct     *last_wakee;

    /*
     * recent_used_cpu is initially set as the last CPU used by a task
     * that wakes affine another task. waker/wakee relationships can
     * push tasks around a CPU where each wakeup moves to the next one.
     * Tracking a recently used CPU allows a quick search for a recently
     * used CPU that may be idle.
     */
    int                   recent_used_cpu;
    int                   wake_cpu;
#endif
    int                   on_rq;

    int                   prio;
    int                   static_prio;
    int                   normal_prio;
    unsigned int          rt_priority;

    struct sched_entity    se;
    struct sched_rt_entity rt;
```

```

    struct sched_dl_entity    dl;
    const struct sched_class  *sched_class;

#ifdef CONFIG_SCHED_CORE
    struct rb_node            core_node;
    unsigned long             core_cookie;
    unsigned int              core_occupation;
#endif

#ifdef CONFIG_CGROUP_SCHED
    struct task_group         *sched_task_group;
#endif

#ifdef CONFIG_UCLAMP_TASK
    /*
     * Clamp values requested for a scheduling entity.
     * Must be updated with task_rq_lock() held.
     */
    struct uclamp_se          uclamp_req[UCLAMP_CNT];
    /*
     * Effective clamp values used for a scheduling entity.
     * Must be updated with task_rq_lock() held.
     */
    struct uclamp_se          uclamp[UCLAMP_CNT];
#endif

    struct sched_statistics    stats;

#ifdef CONFIG_PREEMPT_NOTIFIERS
    /* List of struct preempt_notifier: */
    struct hlist_head          preempt_notifiers;
#endif

#ifdef CONFIG_BLK_DEV_IO_TRACE
    unsigned int               btrace_seq;
#endif

    unsigned int               policy;
    int                         nr_cpus_allowed;
    const cpumask_t            *cpus_ptr;
    cpumask_t                  *user_cpus_ptr;
    cpumask_t                  cpus_mask;
    void                        *migration_pending;
#ifdef CONFIG_SMP
    unsigned short              migration_disabled;
#endif
    unsigned short              migration_flags;

#ifdef CONFIG_PREEMPT_RCU
    int                         rcu_read_lock_nesting;
    union rcu_special           rcu_read_unlock_special;
    struct list_head            rcu_node_entry;
    struct rcu_node             *rcu_blocked_node;
#endif /* #ifdef CONFIG_PREEMPT_RCU */

#ifdef CONFIG_TASKS_RCU
    unsigned long               rcu_tasks_nvcsw;
    u8                          rcu_tasks_holdout;

```

```

    u8            rcu_tasks_idx;
    int           rcu_tasks_idle_cpu;
    struct list_head rcu_tasks_holdout_list;
#endif /* #ifdef CONFIG_TASKS_RCU */

#ifdef CONFIG_TASKS_TRACE_RCU
    int           trc_reader_nesting;
    int           trc_ipi_to_cpu;
    union rcu_special trc_reader_special;
    bool          trc_reader_checked;
    struct list_head trc_holdout_list;
#endif /* #ifdef CONFIG_TASKS_TRACE_RCU */

    struct sched_info sched_info;

    struct list_head tasks;
#ifdef CONFIG_SMP
    struct plist_node pushable_tasks;
    struct rb_node pushable_dl_tasks;
#endif

    struct mm_struct *mm;
    struct mm_struct *active_mm;

    /* Per-thread vma caching: */
    struct vmacache vmacache;

#ifdef SPLIT_RSS_COUNTING
    struct task_rss_stat rss_stat;
#endif
    int exit_state;
    int exit_code;
    int exit_signal;
    /* The signal sent when the parent dies: */
    int pdeath_signal;
    /* JOBCTL_*, siglock protected: */
    unsigned long jobctl;

    /* Used for emulating ABI behavior of previous Linux versions: */
    unsigned int personality;

    /* Scheduler bits, serialized by scheduler locks: */
    unsigned sched_reset_on_fork:1;
    unsigned sched_contributes_to_load:1;
    unsigned sched_migrated:1;
#ifdef CONFIG_PSI
    unsigned sched_psi_wake_requeue:1;
#endif

    /* Force alignment to the next boundary: */
    unsigned :0;

    /* Unserialized, strictly 'current' */

    /*
     * This field must not be in the scheduler word above due to wakelist
     * queueing no longer being serialized by p->on_cpu. However:
     */

```

```

* p->XXX = X;          ttwu()
* schedule()           if (p->on_rq && ..) // false
* smp_mb__after_spinlock(); if (smp_load_acquire(&p->on_cpu) && //true
* deactivate_task()      ttwu_queue_wake_list())
* p->on_rq = 0;          p->sched_remote_wakeup = Y;
*
* guarantees all stores of 'current' are visible before
* ->sched_remote_wakeup gets used, so it can be in this word.
*/
unsigned          sched_remote_wakeup:1;

/* Bit to tell LSMs we're in execve(): */
unsigned          in_execve:1;
unsigned          in_iowait:1;
#ifdef TIF_RESTORE_SIGMASK
unsigned          restore_sigmask:1;
#endif
#ifdef CONFIG_MEMCG
unsigned          in_user_fault:1;
#endif
#ifdef CONFIG_COMPAT_BRK
unsigned          brk_randomized:1;
#endif
#ifdef CONFIG_CGROUPS
/* disallow userland-initiated cgroup migration */
unsigned          no_cgroup_migration:1;
/* task is frozen/stopped (used by the cgroup freezer) */
unsigned          frozen:1;
#endif
#ifdef CONFIG_BLK_CGROUP
unsigned          use_memdelay:1;
#endif
#ifdef CONFIG_PSI
/* Stalled due to lack of memory */
unsigned          in_memstall:1;
#endif
#ifdef CONFIG_PAGE_OWNER
/* Used by page_owner=on to detect recursion in page tracking. */
unsigned          in_page_owner:1;
#endif
#ifdef CONFIG_EVENTFD
/* Recursion prevention for eventfd_signal() */
unsigned          in_eventfd_signal:1;
#endif
#ifdef CONFIG_IOMMU_SVA
unsigned          pasid_activated:1;
#endif
#ifdef CONFIG_CPU_SUP_INTEL
unsigned          reported_split_lock:1;
#endif

unsigned long      atomic_flags; /* Flags requiring atomic access. */

struct restart_block restart_block;

pid_t              pid;
pid_t              tgid;

```

```

#ifdef CONFIG_STACKPROTECTOR
    /* Canary value for the -fstack-protector GCC feature: */
    unsigned long        stack_canary;
#endif

/*
 * Pointers to the (original) parent process, youngest child, younger
sibling,
 * older sibling, respectively. (p->father can be replaced with
 * p->real_parent->pid)
 */

/* Real parent process: */
struct task_struct __rcu    *real_parent;

/* Recipient of SIGCHLD, wait4() reports: */
struct task_struct __rcu    *parent;

/*
 * Children/sibling form the list of natural children:
 */
struct list_head    children;
struct list_head    sibling;
struct task_struct    *group_leader;

/*
 * 'ptraced' is the list of tasks this task is using ptrace() on.
 *
 * This includes both natural children and PTTRACE_ATTACH targets.
 * 'ptrace_entry' is this task's link on the p->parent->ptraced list.
 */
struct list_head    ptraced;
struct list_head    ptrace_entry;

/* PID/PID hash table linkage. */
struct pid    *thread_pid;
struct hlist_node    pid_links[PIDTYPE_MAX];
struct list_head    thread_group;
struct list_head    thread_node;

struct completion    *vfork_done;

/* CLONE_CHILD_SETTID: */
int __user    *set_child_tid;

/* CLONE_CHILD_CLEARTID: */
int __user    *clear_child_tid;

/* PF_KTHREAD | PF_IO_WORKER */
void    *worker_private;

u64    utime;
u64    stime;
#ifdef CONFIG_ARCH_HAS_SCALED_CPUTIME
    u64    utimescaled;
    u64    stimescaled;
#endif
u64    gtime;
struct prev_cputime    prev_cputime;

```



```

#ifdef CONFIG_VIRT_CPU_ACCOUNTING_GEN
    struct vtime          vtime;
#endif

#ifdef CONFIG_NO_HZ_FULL
    atomic_t              tick_dep_mask;
#endif

/* Context switch counts: */
unsigned long            nvcs;
unsigned long            nivcs;

/* Monotonic time in nsecs: */
u64                      start_time;

/* Boot based time in nsecs: */
u64                      start_boottime;

/* MM fault and swap info: this can arguably be seen as either mm-specific or
thread-specific: */
unsigned long            min_flt;
unsigned long            maj_flt;

/* Empty if CONFIG_POSIX_CPUTIMERS=n */
struct posix_cputimers    posix_cputimers;

#ifdef CONFIG_POSIX_CPU_TIMERS_TASK_WORK
    struct posix_cputimers_work posix_cputimers_work;
#endif

/* Process credentials: */

/* Tracer's credentials at attach: */
const struct cred __rcu    *ptracer_cred;

/* Objective and real subjective task credentials (COW): */
const struct cred __rcu    *real_cred;

/* Effective (overridable) subjective task credentials (COW): */
const struct cred __rcu    *cred;

#ifdef CONFIG_KEYS
    /* Cached requested key. */
    struct key            *cached_requested_key;
#endif

/*
 * executable name, excluding path.
 *
 * - normally initialized setup_new_exec()
 * - access it with [gs]et_task_comm()
 * - lock it with task_lock()
 */
char                      comm[TASK_COMM_LEN];

    struct nameidata      *nameidata;

#ifdef CONFIG_SYSVIPC
    struct sysv_sem        sysvsem;

```

```

    struct sysv_shm          sysvshm;
#endif
#ifdef CONFIG_DETECT_HUNG_TASK
    unsigned long            last_switch_count;
    unsigned long            last_switch_time;
#endif
    /* Filesystem information: */
    struct fs_struct          *fs;

    /* Open file information: */
    struct files_struct       *files;

#ifdef CONFIG_IO_URING
    struct io_uring_task      *io_uring;
#endif

    /* Namespaces: */
    struct nsproxy            *nsproxy;

    /* Signal handlers: */
    struct signal_struct      *signal;
    struct sighand_struct     __rcu    *sighand;
    sigset_t                  blocked;
    sigset_t                  real_blocked;
    /* Restored if set_restore_sigmask() was used: */
    sigset_t                  saved_sigmask;
    struct sigpending         pending;
    unsigned long             sas_ss_sp;
    size_t                    sas_ss_size;
    unsigned int              sas_ss_flags;

    struct callback_head      *task_works;

#ifdef CONFIG_AUDIT
#ifdef CONFIG_AUDITSYSCALL
    struct audit_context      *audit_context;
#endif
#endif
    kuid_t                    loginuid;
    unsigned int              sessionid;
#endif
    struct seccomp            seccomp;
    struct syscall_user_dispatch syscall_dispatch;

    /* Thread group tracking: */
    u64                       parent_exec_id;
    u64                       self_exec_id;

    /* Protection against (de-)allocation: mm, files, fs, tty, keyrings,
    mems_allowed, mempolicy: */
    spinlock_t                alloc_lock;

    /* Protection of the PI data structures: */
    raw_spinlock_t            pi_lock;

    struct wake_q_node        wake_q;

#ifdef CONFIG_RT_MUTEXES
    /* PI waiters blocked on a rt_mutex held by this task: */

```

```

    struct rb_root_cached    pi_waiters;
    /* Updated under owner's pi_lock and rq lock */
    struct task_struct       *pi_top_task;
    /* Deadlock detection and priority inheritance handling: */
    struct rt_mutex_waiter    *pi_blocked_on;
#endif

#ifdef CONFIG_DEBUG_MUTEXES
    /* Mutex deadlock detection: */
    struct mutex_waiter       *blocked_on;
#endif

#ifdef CONFIG_DEBUG_ATOMIC_SLEEP
    int                        non_block_count;
#endif

#ifdef CONFIG_TRACE_IRQFLAGS
    struct irqtrace_events     irqtrace;
    unsigned int               hardirq_threaded;
    u64                         hardirq_chain_key;
    int                         softirqs_enabled;
    int                         softirq_context;
    int                         irq_config;
#endif

#ifdef CONFIG_PREEMPT_RT
    int                         softirq_disable_cnt;
#endif

#ifdef CONFIG_LOCKDEP
#define MAX_LOCK_DEPTH        48UL
    u64                         curr_chain_key;
    int                         lockdep_depth;
    unsigned int                lockdep_recursion;
    struct held_lock             held_locks[MAX_LOCK_DEPTH];
#endif

#ifdef CONFIG_UBSAN && !defined(CONFIG_UBSAN_TRAP)
    unsigned int                in_ubsan;
#endif

    /* Journaling filesystem info: */
    void                        *journal_info;

    /* Stacked block device info: */
    struct bio_list              *bio_list;

    /* Stack plugging: */
    struct blk_plug              *plug;

    /* VM state: */
    struct reclaim_state         *reclaim_state;

    struct backing_dev_info      *backing_dev_info;

    struct io_context            *io_context;

#ifdef CONFIG_COMPACTION
    struct capture_control       *capture_control;

```

```

#endif
/* Ptrace state: */
unsigned long      ptrace_message;
kernel_siginfo_t   *last_siginfo;

struct task_io_accounting ioac;
#ifdef CONFIG_PSI
/* Pressure stall state */
unsigned int       psi_flags;
#endif
#ifdef CONFIG_TASK_XACCT
/* Accumulated RSS usage: */
u64                acct_rss_mem1;
/* Accumulated virtual memory usage: */
u64                acct_vm_mem1;
/* stime + utime since last update: */
u64                acct_timexpd;
#endif
#ifdef CONFIG_CPUSETS
/* Protected by ->alloc_lock: */
nodemask_t         mems_allowed;
/* Sequence number to catch updates: */
seqcount_spinlock_t mems_allowed_seq;
int                cpuset_mem_spread_rotor;
int                cpuset_slab_spread_rotor;
#endif
#ifdef CONFIG_CGROUPS
/* Control Group info protected by css_set_lock: */
struct css_set __rcu *cgroups;
/* cg_list protected by css_set_lock and tsk->alloc_lock: */
struct list_head   cg_list;
#endif
#ifdef CONFIG_X86_CPU_RESCTRL
u32                closid;
u32                rmid;
#endif
#ifdef CONFIG_FUTEX
struct robust_list_head __user *robust_list;
#endif
#ifdef CONFIG_COMPAT
struct compat_robust_list_head __user *compat_robust_list;
#endif
struct list_head     pi_state_list;
struct futex_pi_state *pi_state_cache;
struct mutex         futex_exit_mutex;
unsigned int         futex_state;
#endif
#ifdef CONFIG_PERF_EVENTS
struct perf_event_context *perf_event_ctxp[perf_nr_task_contexts];
struct mutex         perf_event_mutex;
struct list_head     perf_event_list;
#endif
#ifdef CONFIG_DEBUG_PREEMPT
unsigned long        preempt_disable_ip;
#endif
#ifdef CONFIG_NUMA
/* Protected by alloc_lock: */
struct mempolicy     *mempolicy;
short                il_prev;

```

```

        short                pref_node_fork;
#endif
#ifdef CONFIG_NUMA_BALANCING
    int                numa_scan_seq;
    unsigned int        numa_scan_period;
    unsigned int        numa_scan_period_max;
    int                numa_preferred_nid;
    unsigned long        numa_migrate_retry;
    /* Migration stamp: */
    u64                node_stamp;
    u64                last_task_numa_placement;
    u64                last_sum_exec_runtime;
    struct callback_head    numa_work;

    /*
     * This pointer is only modified for current in syscall and
     * pagefault context (and for tasks being destroyed), so it can be read
     * from any of the following contexts:
     * - RCU read-side critical section
     * - current->numa_group from everywhere
     * - task's runqueue locked, task not running
     */
    struct numa_group __rcu    *numa_group;

    /*
     * numa_faults is an array split into four regions:
     * faults_memory, faults_cpu, faults_memory_buffer, faults_cpu_buffer
     * in this precise order.
     *
     * faults_memory: Exponential decaying average of faults on a per-node
     * basis. Scheduling placement decisions are made based on these
     * counts. The values remain static for the duration of a PTE scan.
     * faults_cpu: Track the nodes the process was running on when a NUMA
     * hinting fault was incurred.
     * faults_memory_buffer and faults_cpu_buffer: Record faults per node
     * during the current scan window. When the scan completes, the counts
     * in faults_memory and faults_cpu decay and these values are copied.
     */
    unsigned long        *numa_faults;
    unsigned long        total_numa_faults;

    /*
     * numa_faults_locality tracks if faults recorded during the last
     * scan window were remote/local or failed to migrate. The task scan
     * period is adapted based on the locality of the faults with different
     * weights depending on whether they were shared or private faults
     */
    unsigned long        numa_faults_locality[3];

    unsigned long        numa_pages_migrated;
#endif /* CONFIG_NUMA_BALANCING */

#ifdef CONFIG_RSEQ
    struct rseq __user *rseq;
    u32 rseq_sig;
    /*
     * RmW on rseq_event_mask must be performed atomically
     * with respect to preemption.

```

```

    */
    unsigned long rseq_event_mask;
#endif

    struct tlbflush_unmap_batch tlb_ubic;

    union {
        refcount_t      rcu_users;
        struct rcu_head  rcu;
    };

    /* Cache last used pipe for splice(): */
    struct pipe_inode_info      *splice_pipe;

    struct page_frag          task_frag;

#ifdef CONFIG_TASK_DELAY_ACCT
    struct task_delay_info      *delays;
#endif

#ifdef CONFIG_FAULT_INJECTION
    int          make_it_fail;
    unsigned int fail_nth;
#endif

    /*
     * When (nr_dirtied >= nr_dirtied_pause), it's time to call
     * balance_dirty_pages() for a dirty throttling pause:
     */
    int          nr_dirtied;
    int          nr_dirtied_pause;
    /* Start of a write-and-pause period: */
    unsigned long dirty_paused_when;

#ifdef CONFIG_LATENCYTOP
    int          latency_record_count;
    struct latency_record latency_record[LT_SAVECOUNT];
#endif

    /*
     * Time slack values; these are used to round up poll() and
     * select() etc timeout values. These are in nanoseconds.
     */
    u64          timer_slack_ns;
    u64          default_timer_slack_ns;

#ifdef CONFIG_KASAN_GENERIC || defined(CONFIG_KASAN_SW_TAGS)
    unsigned int kasan_depth;
#endif

#ifdef CONFIG_KCSAN
    struct kcsan_ctx kcsan_ctx;
#endif
#ifdef CONFIG_TRACE_IRQFLAGS
    struct irqtrace_events kcsan_save_irqtrace;
#endif
#ifdef CONFIG_KCSAN_WEAK_MEMORY
    int          kcsan_stack_depth;
#endif
#endif

```

```

#if IS_ENABLED(CONFIG_KUNIT)
    struct kunit          *kunit_test;
#endif

#ifdef CONFIG_FUNCTION_GRAPH_TRACER
    /* Index of current stored address in ret_stack: */
    int          curr_ret_stack;
    int          curr_ret_depth;

    /* Stack of return addresses for return function tracing: */
    struct ftrace_ret_stack    *ret_stack;

    /* Timestamp for last schedule: */
    unsigned long long    ftrace_timestamp;

    /*
     * Number of functions that haven't been traced
     * because of depth overrun:
     */
    atomic_t          trace_overrun;

    /* Pause tracing: */
    atomic_t          tracing_graph_pause;
#endif

#ifdef CONFIG_TRACING
    /* State flags for use by tracers: */
    unsigned long    trace;

    /* Bitmask and counter of trace recursion: */
    unsigned long    trace_recursion;
#endif /* CONFIG_TRACING */

#ifdef CONFIG_KCOV
    /* See kernel/kcov.c for more details. */

    /* Coverage collection mode enabled for this task (0 if disabled): */
    unsigned int    kcov_mode;

    /* Size of the kcov_area: */
    unsigned int    kcov_size;

    /* Buffer for coverage collection: */
    void          *kcov_area;

    /* KCOV descriptor wired with this task or NULL: */
    struct kcov    *kcov;

    /* KCOV common handle for remote coverage collection: */
    u64          kcov_handle;

    /* KCOV sequence number: */
    int          kcov_sequence;

    /* Collect coverage from softirq context: */
    unsigned int    kcov_softirq;
#endif

```

```

#ifdef CONFIG_MEMCG
    struct mem_cgroup      *memcg_in_oom;
    gfp_t                  memcg_oom_gfp_mask;
    int                    memcg_oom_order;

    /* Number of pages to reclaim on returning to userland: */
    unsigned int           memcg_nr_pages_over_high;

    /* Used by memcontrol for targeted memcg charge: */
    struct mem_cgroup      *active_memcg;
#endif

#ifdef CONFIG_BLK_CGROUP
    struct request_queue    *throttle_queue;
#endif

#ifdef CONFIG_UPROBES
    struct uprobe_task      *utask;
#endif
#if defined(CONFIG_BCACHE) || defined(CONFIG_BCACHE_MODULE)
    unsigned int            sequential_io;
    unsigned int            sequential_io_avg;
#endif
    struct kmap_ctrl        kmap_ctrl;
#ifdef CONFIG_DEBUG_ATOMIC_SLEEP
    unsigned long           task_state_change;
#endif
#ifdef CONFIG_PREEMPT_RT
    unsigned long           saved_state_change;
#endif
#endif
    int                    pagefault_disabled;
#ifdef CONFIG_MMU
    struct task_struct      *oom_reaper_list;
    struct timer_list        oom_reaper_timer;
#endif
#ifdef CONFIG_VMAP_STACK
    struct vm_struct        *stack_vm_area;
#endif
#ifdef CONFIG_THREAD_INFO_IN_TASK
    /* A live task holds one reference: */
    refcount_t              stack_refcount;
#endif
#ifdef CONFIG_LIVEPATCH
    int patch_state;
#endif
#ifdef CONFIG_SECURITY
    /* Used by LSM modules for access restriction: */
    void                    *security;
#endif
#ifdef CONFIG_BPF_SYSCALL
    /* Used by BPF task local storage */
    struct bpf_local_storage __rcu *bpf_storage;
    /* Used for BPF run context */
    struct bpf_run_ctx      *bpf_ctx;
#endif

#ifdef CONFIG_GCC_PLUGIN_STACKLEAK
    unsigned long           lowest_stack;

```



```

        unsigned long          prev_lowest_stack;
#endif

#ifdef CONFIG_X86_MCE
    void __user          *mce_vaddr;
    __u64                mce_kflags;
    u64                  mce_addr;
    __u64                mce_ripv : 1,
                        mce_whole_page : 1,
                        __mce_reserved : 62;
    struct callback_head  mce_kill_me;
    int                  mce_count;
#endif

#ifdef CONFIG_KRETPROBES
    struct llist_head     kretprobe_instances;
#endif
#ifdef CONFIG_RETHOOK
    struct llist_head     rethooks;
#endif

#ifdef CONFIG_ARCH_HAS_PARANOID_L1D_FLUSH
    /*
     * If L1D flush is supported on mm context switch
     * then we use this callback head to queue kill work
     * to kill tasks that are not running on SMT disabled
     * cores
     */
    struct callback_head  l1d_flush_kill;
#endif

    /*
     * New fields for task_struct should be added above here, so that
     * they are included in the randomized portion of task_struct.
     */
    randomized_struct_fields_end

    /* CPU-specific state of this task: */
    struct thread_struct   thread;

    /*
     * WARNING: on x86, 'thread_struct' contains a variable-sized
     * structure. It *MUST* be at the end of 'task_struct'.
     *
     * Do not put anything below here!
     */
};

```

Add new system call to list all processes

需要列出每个进程的名字(comm)、进程ID号(pid)、父进程ID号、进程状态、学号姓名等

进程结构体: struct task_struct *p;

双向链表:

```
struct list_head {
    struct list_head *next, *prev;
};
```

第一个进程:init_task

当前的集成:current

进程的兄弟进程:struct list_head tasks;

进程p的下一个进程:next_task(p)

```
#define next_task(p) \
    list_entry_rcu((p)->tasks.next, struct task_struct, tasks)
```

遍历所有进程:

```
#define for_each_process(p) \
    for (p = &init_task ; (p = next_task(p)) != &init_task ; )
```

进程的状态:static inline char task_state_to_char(struct task_struct *tsk)

实现框架

kernel/sys.c

在文件SYSCALL_DEFINE0(gettid)函数之后，添加如下行

```
SYSCALL_DEFINE0(schello)
{
    struct task_struct *p;
    printf("Hello new system call schello!\n");
    printk("%-20s %-6s %-6s\n", "Name", "Pid", "Stat");
    ...

    return 0;
}
```

End.